Results: [20, 35] Crossings

Jake Weatherhead

The CNN (88M params) and ViT (86M params) were fine-tuned on a new dataset $\mathcal{D}$ comprising unknot $\mathcal{U}$ and non-trivial knot $\mathcal{K}$ diagrams partitioned into three disjoint data splits $S := \{\text{train, val, test}\}$, where:

$$\mathcal{D} \;=\; \biguplus_{s \in S} \mathcal{D}_s$$

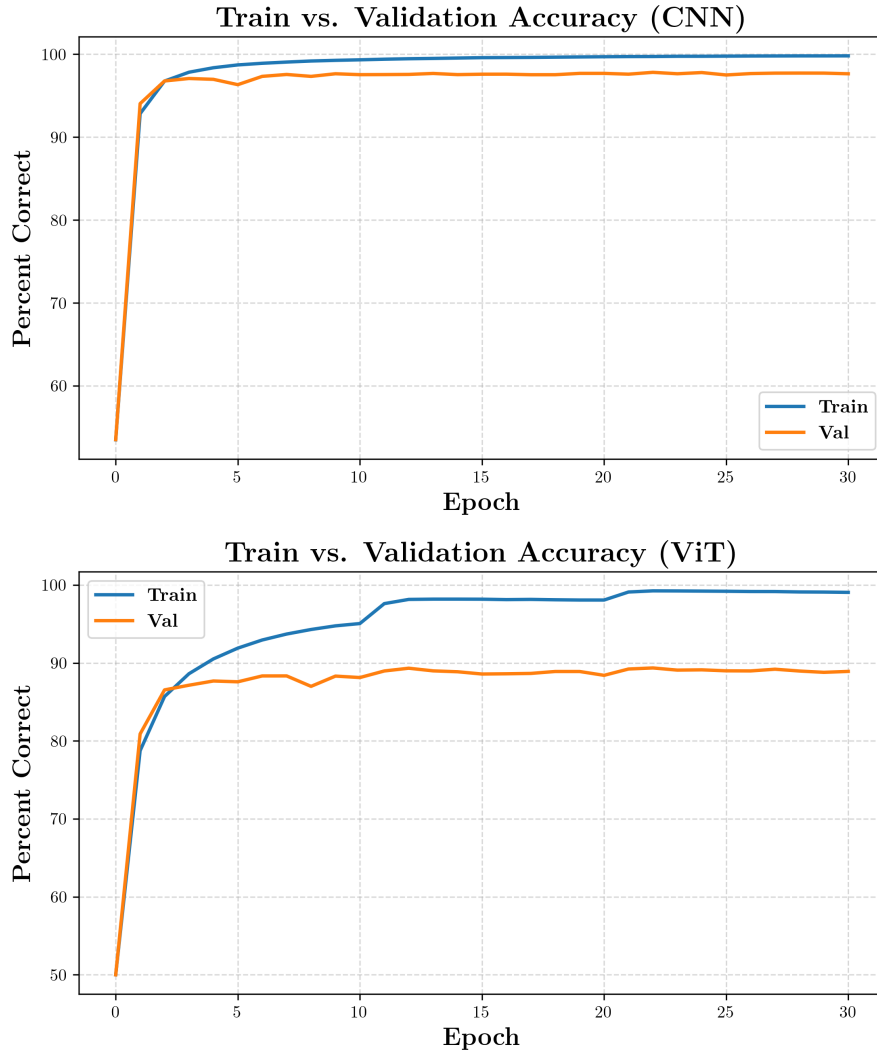$$|\mathcal{D}| \;=\; 560{,}000 \;\; \text{diagrams},$$

$$|\mathcal{D}_{\text{train}}| \;=\; 0.8 \cdot |\mathcal{D}| \;=\; 448{,}000 \;\text{diagrams},$$

$$|\mathcal{D}_{\text{val}}| \;=\; |\mathcal{D}_{\text{test}}| \;=\; 0.1 \cdot |\mathcal{D}| \;=\; 56{,}000 \;\text{diagrams},$$

$$\text{where} \;\; \forall s \in S, \;\; s \in \mathcal{U}_s \cup \mathcal{K}_s \;\; \text{and} \;\; |\mathcal{U}_s| = |\mathcal{K}_s|.$$

Each split contained knot diagrams for every crossing count $n \in N := \{20, 21, \dots, 35\}$. For all $x, y \in N$, $x \neq y$, each split in S contained an equal number of x and y-crossing unknot and non-trivial knot diagrams.

Learning Curves

Test Accuracy by Crossing Count

Let CNN* and ViT* denote the models with the highest test classification accuracies.

CNN* Saliency Maps

A sample of CNN* saliency maps for knot diagrams in $\mathcal{D}_{\text{test}}$.

### True Positives: Unknots Predicted as Unknots

Figure 1: Of the 28,000 unknot diagrams in $\mathcal{D}_{test}$, CNN* produced 27,677 true positives.

False Negatives: Unknots Predicted as Non-Trivial

Figure 2: Of the 28,000 unknot diagrams in $\mathcal{D}_{test}$, CNN* produced 323 false negatives.

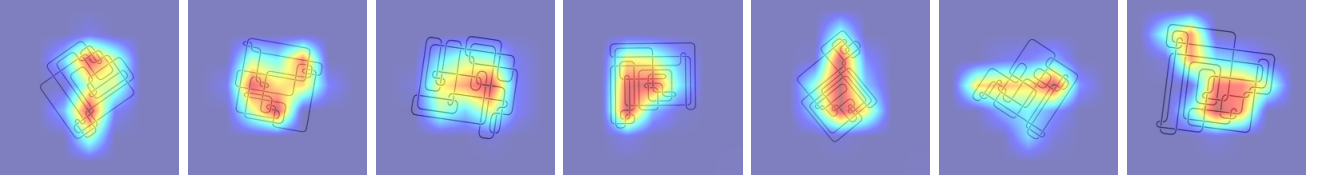**True Negatives: Non-Trivials Predicted as Non-Trivial**

Figure 3: Of the 28,000 non-trivial knot diagrams in $\mathcal{D}_{test}$, CNN* produced 27,020 true negatives.

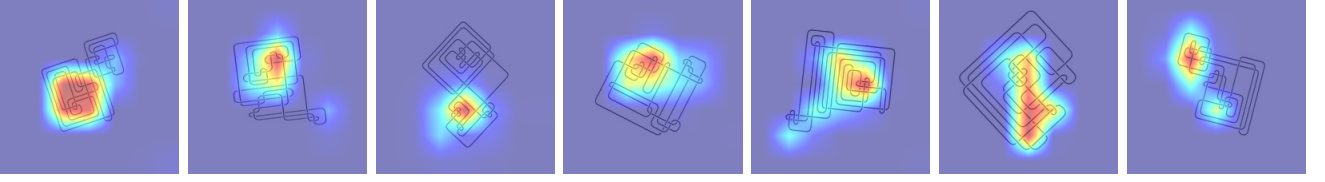**False Positives: Non-Trivials Predicted as Unknots**

Figure 4: Of the 28,000 non-trivial knot diagrams in $\mathcal{D}_{test}$, CNN* produced 980 false positives.

ViT* Saliency Maps

A sample of ViT* saliency maps for knot diagrams in $\mathcal{D}_{test}$.

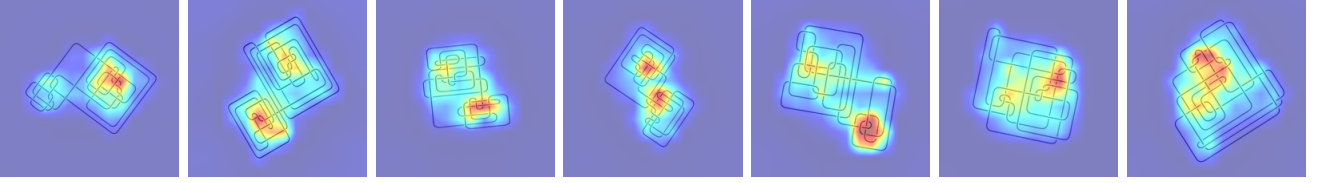**True Positives: Unknots Predicted as Unknots**

Figure 5: Of the 28,000 unknot diagrams in $\mathcal{D}_{test}$, ViT* produced 25,630 true positives.

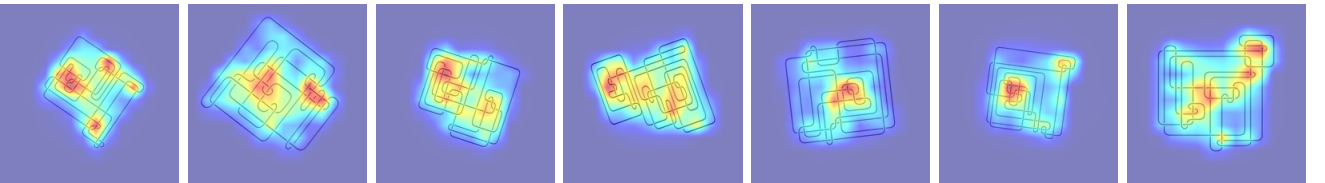**False Negatives: Unknots Predicted as Non-Trivial**

Figure 6: Of the 28,000 unknot diagrams in $\mathcal{D}_{test}$, ViT* produced 2,370 false negatives.

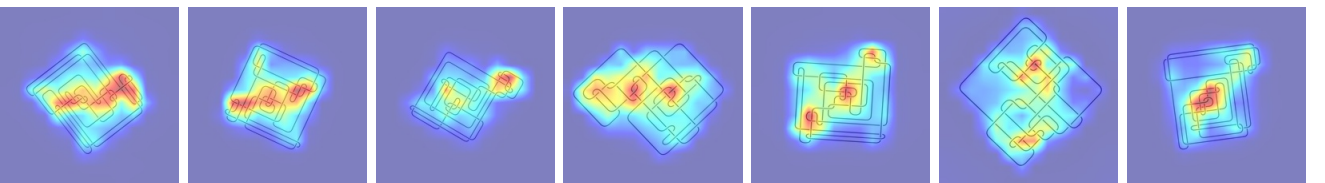**True Negatives: Non-Trivials Predicted as Non-Trivial**

Figure 7: Of the 28,000 non-trivial knot diagrams in $\mathcal{D}_{test}$, ViT* produced 24,252 true negatives.

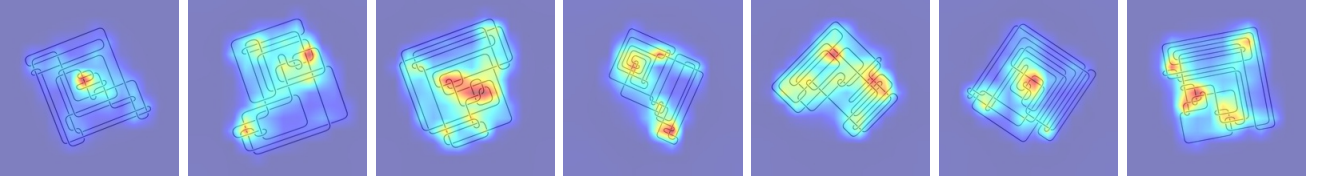## False Positives: Non-Trivials Predicted as Unknots

Figure 8: Of the 28,000 non-trivial knot diagrams in $\mathcal{D}_{test}$, ViT* produced 3,748 false positives.

Further Evidence of Learning

Let m_v denote checkpoint model $m \in \{\text{CNN}, \text{ViT}\}$ with validation accuracy v. The figures below are examples of a model m initially misclassifying a knot diagram $d \in \mathcal{D}_{test}$, later learning to correctly classify d. Let p and t respectively denote the predicted and true triviality of d as given by model m_v.