

Results: [20, 35] Crossings

Jake Weatherhead

November 8, 2025

The CNN (88M params) and ViT (86M params) were fine-tuned on a new dataset \mathcal{D} comprising unknot \mathcal{U} and non-trivial knot \mathcal{K} diagrams partitioned into three disjoint data splits $S := \{\text{train}, \text{val}, \text{test}\}$, where:

$$\mathcal{D} = \bigcup_{s \in S} \mathcal{D}_s$$

$$|\mathcal{D}| = 560,000 \text{ diagrams,}$$

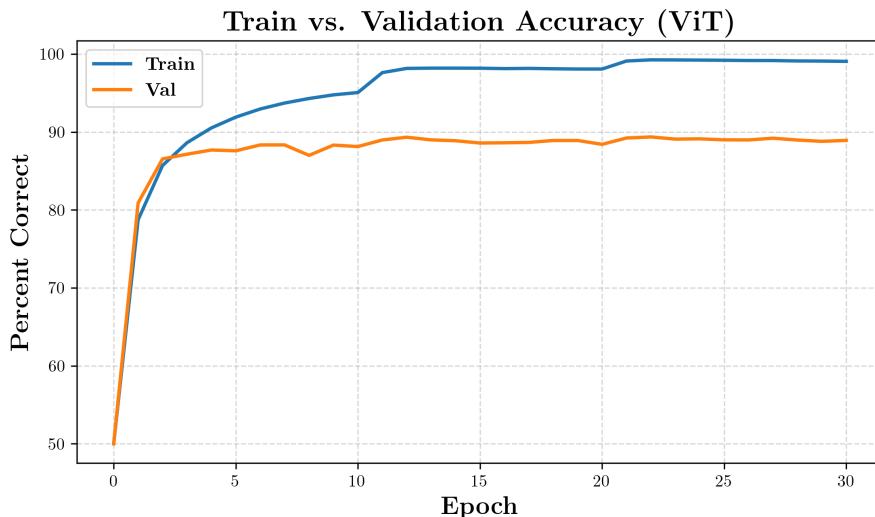
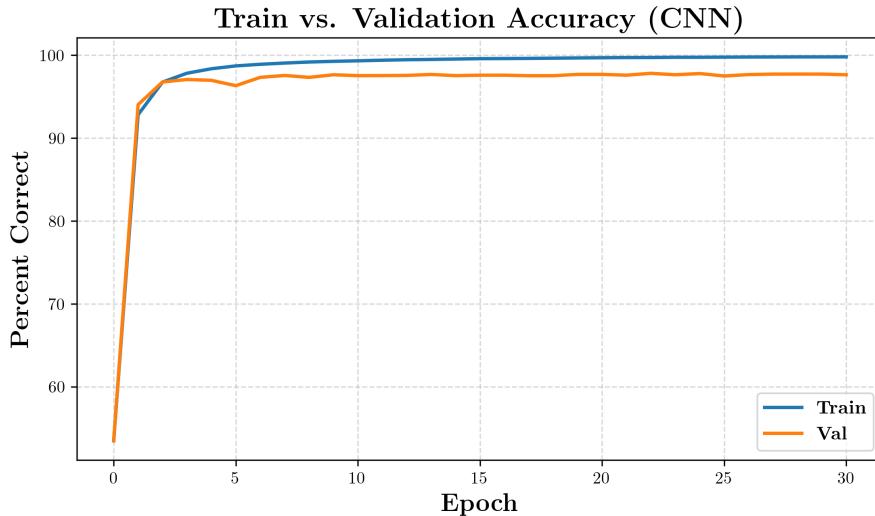
$$|\mathcal{D}_{\text{train}}| = 0.8 \cdot |\mathcal{D}| = 448,000 \text{ diagrams,}$$

$$|\mathcal{D}_{\text{val}}| = |\mathcal{D}_{\text{test}}| = 0.1 \cdot |\mathcal{D}| = 56,000 \text{ diagrams,}$$

$$\text{where } \forall s \in S, s \in \mathcal{U}_s \cup \mathcal{K}_s \text{ and } |\mathcal{U}_s| = |\mathcal{K}_s|.$$

Each split contained knot diagrams for every crossing count $n \in N := \{20, 21, \dots, 35\}$. For all distinct $p, q \in N$, each split contained an equal number of p and q -crossing unknot and non-trivial knot diagrams.

Learning Curves



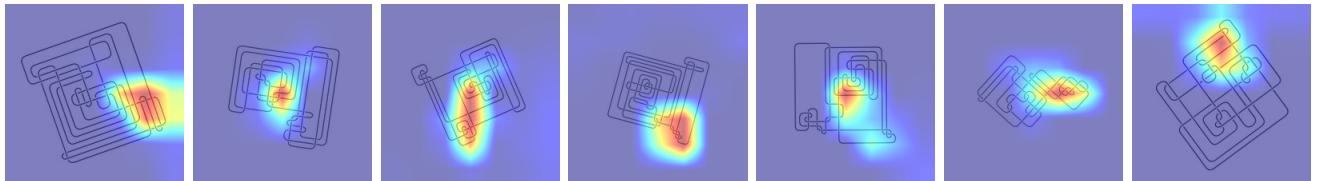
Test Accuracy by Crossing Count

CNN* Saliency Maps

Saliency maps for knot diagrams in $\mathcal{D}_{\text{test}}$ for the best-performing CNN (CNN*).

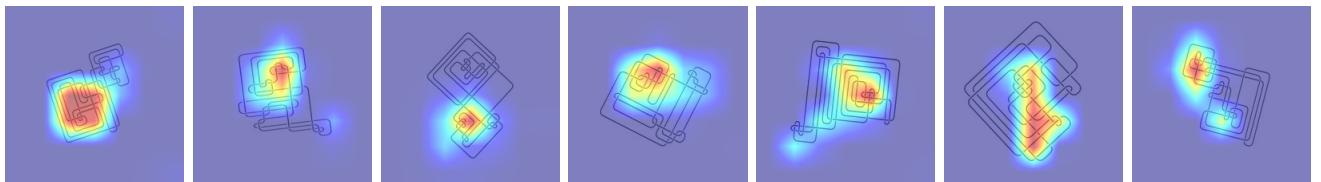
CNN* True Positives: True Unknots Predicted as Unknots

CNN* False Negatives: True Unknots Predicted as Knots



CNN* True Negatives: True Knots Predicted as Knots

CNN* False Positives: True Knots Predicted as Unknots



ViT* Saliency Maps

Saliency maps for knot diagrams in $\mathcal{D}_{\text{test}}$ for the best-performing ViT (ViT*).

ViT TP: True Unknots Predicted as Unknots

ViT FN: True Unknots Predicted as Knots

ViT TN: True Knots Predicted as Knots

ViT FP: True Knots Predicted as Unknots